

EC2 Pricing Study Notes

Andrew Brown

On Demand

When you launch an EC2 instance it is by default On-Demand Pricing

Key Points

- On-Demand has **no up-front payment** and **no long-term commitment**
- You are charged by the **hour** or by the **minute** (varies based on EC2 Instance Types)
- On-Demand is for applications where the workload is for **short-term, spikey,** or **unpredictable.**
- Used for when you have a **new app** for development or you want to run an experiment

EC2- Reserved Instances (RI)

Reserved Instance Pricing: Designed for applications that have a **steady-state, predictable usage**, or require **reserved capacity**.

Key Points

- Reduced pricing is based on **Term x Class Offering x Payment Option**
- **Types of payment options:**

- **Standard:** Up to **75%** reduced pricing compared to on-demand. Cannot change RI Attributes
- **Convertible:** Up to **54%** reduced pricing compared to on-demand. Allows you to change RI Attributes if greater or equal in value.
- **Scheduled:** You reserve instances for specific time periods. **Savings vary.**
- **Terms-** You commit to a 1 year or 3 year contract. The longer the term the greater savings
- **Payment Options-** All Upfront, Partial Upfront, and No Upfront. The greater upfront the greater the savings.
- **RIs** can be shared between multiple accounts within an organization
- **Unused RIs** can be sold in the **Reserved Instance Marketplace.**

Spot Instances

Amazon EC2 Spot Instances: Run fault-tolerant workloads for up to **90% off**

Key Points

- Amazon EC2 Spot Instances let you take advantage of **unused EC2 capacity** in the AWS cloud.
- Designed for applications that have flexible start and end times or applications that are only feasible at **very low** compute costs.
- AWS Batch is an easy and convenient way to use Spot Pricing
- **Termination Conditions:**
 - Instances can be terminated by AWS **at anytime.**
 - If your instance is terminated by AWS, **you don't get charged** for a partial hour of usage.

- If you terminate an instance **you will still be charged** for an hour that it ran.

Dedicated Host Instances

Dedicated Host Instances: dedicated hardware to support existing software licenses and improve compliance.

Key Points

- Most expensive
- Designed to meet regulatory requirements.
- Use when you have strict **server-bound licensing** that won't support multi-tenancy or cloud deployments
- **Multi-tenant**: When multiple customers are running workloads on the same hardware. **Virtual Isolation** is what separate customers (Like a multi-tenant apartment)
- **Single Tenant**: When a single customer has dedicated hardware. **Physical Isolation** is what separates customers(Like a single-family home)
- Offered in both **On-demand** and **Reserved** (70% of on-demand pricing)
- **Enterprises** and **Large Organizations** may have security concerns or obligations against sharing the same hardware with other AWS Customers

Savings Plans

AWS Savings Plans was introduced as an alternative to Reserved Instance (RI) Pricing to simplify purchasing. They sound similar but there are some key differences.

With **Savings Plans** you have to commit up-front to:

- 1 year contract
- 3 years contract

and any usage over the committed usage is charged at On-Demand rate

With **Reserved Instances**, you can commit to full, partial or up-front. You can commit to

- 1 year contract
- 3 years contract

With **Savings Plans** you can only save with EC2 and Fargate

With **Reserved Instances** you can only save with EC2, RDS, Redshift and ElastiCache

With *Savings Plans* **you commit to a minimum dollar per hour spend** With *Reserved Instances* you commit to a specific amount of utilization

With **Convertible Reserved Instance** you can resell your utilized resources in the AWS Marketplace With **Savings Plans** you cannot resell in the AWS Marketplace.

Cheat Sheet

■ **EC2 has 4 pricing models**

- 1 ■ On-Demand
- 2 ■ Spot
- 3 ■ Reserved Instances (RI)
- 4 ■ Dedicated

- **On-Demand** (Least commitment)
 - Low cost and flexible
 - Only pay per hour
 - Use Case: short-term, spiky, unpredictable workloads, and first time apps
 - Ideal when your workloads cannot be interrupted.
- **Reserved Instances**- up to 75% off (Best long-term value)
 - **Use case:** steady state or predictable usage
 - Can resell unused reserved instances (Reserved Instance Marketplace)
 - Reduced Pricing is based on **Term x Class Offering x Payment Option**
 - **Payment Terms:** 1 year or 3 years
 - **Payment Options:** All Upfront, Partial Upfront, No Upfront
 - **Class Offerings:**
 - **Standard:** Up to 75% reduced pricing compared to on-demand. Cannot change RI attributes
 - **Convertible:** Up to 54% reduced pricing compared to on-demand. Allows you to change RI Attributes if greater than or equal in value.
 - **Scheduled:** You reserve instances for specific times periods. Savings vary.
- **Spot Pricing**- up to 90% off (Biggest Savings)
 - Request spare computing capacity
 - Flexible start and end times
 - **Use case:** Can handle interruptions like a server randomly stopping and starting
 - **Use case:** For non-critical background jobs
 - Instances can be terminated by AWS **at anytime**
 - If your instance is **terminated by AWS**, you don't get charged for a partial hour of usage.
 - If **you terminate** an instance you will still be charged for any hour that it ran.
- **Dedicated Hosting** (Most Expensive)
 - Dedicated servers
 - Can be on-demand or reserved (up to 70% off)
 - **Use case:** When you need a guarantee of isolated hardware (enterprise requirements)